

AFSAL ASHRAF

📍 Abu Dhabi, United Arab Emirates ✉ email@afsalashraf.com ☎ +971 - 508852463 📧 in/afsalsa 🌐 https://www.afsalashraf.com

SUMMARY

Cloud Architect & AI Platform Engineer with **14+ years of experience** architecting and delivering enterprise-scale **Azure cloud, AI/ML, and cloud-native platforms**, with additional expertise across **Google Cloud (GCP)** and hybrid environments. Specialized in **Generative AI, Agentic AI, RAG, MLOps, Kubernetes, Platform Engineering, Terraform, and DevSecOps**, with hands-on experience across **Microsoft Foundry, Azure OpenAI, Azure Machine Learning, and Azure AI Services**. Experienced in **NVIDIA AI software technologies**, including **NVIDIA AI Enterprise, NIM, NeMo, Triton Inference Server, and TensorRT-LLM**, supporting scalable AI platform and inference architectures.

EXPERIENCE

AI Platform Lead / Cloud Architect | Abu Dhabi Commercial Bank (ADCB)

January 2026 - Present, Abu Dhabi

- Led the architecture and delivery of an enterprise **Azure AI platform** for ADCB, establishing a secure, scalable foundation for production AI workloads and enterprise use cases.
- Architected an **AI Landing Zone using Azure Citadel**, establishing standardized cloud foundations, security controls, networking, and governance for enterprise AI workloads.
- Designed and implemented AI inference capabilities across the UAE region, integrating **Core42 Compass** and **Azure OpenAI** to support enterprise AI workloads and regional deployment requirements.
- Architected AI solutions using **Microsoft Foundry**, implementing **Agentic AI** and **RAG** workflows with GPT and open-source models to enable knowledge-driven enterprise applications.
- Designed and implemented an **Azure AI Gateway**, providing a centralized integration and control layer for enterprise AI model consumption and application workloads.
- Successfully delivered **5 enterprise AI use cases** leveraging **Microsoft Foundry, Azure OpenAI, Agentic AI, and RAG**, translating business requirements into production-ready AI solutions.
- Developed reusable **Terraform** codebases and standalone infrastructure modules for banking workloads, standardizing and accelerating deployment of Azure services across the AI platform.
- Implemented **DevSecOps pipelines using GitHub Actions**, enabling secure, automated, and repeatable infrastructure and application deployments across AI environments.
- Established a comprehensive AI governance framework, defining policies and controls for data privacy, model bias, and ethical AI usage across 15+ enterprise AI applications.

SDE 4 | Cloud & AI Platform Engineering | Walmart Inc

September 2020 - December 2025, Bengaluru, India

- Architected and led enterprise-scale cloud platform engineering across **Google Cloud, Azure, and private cloud**, leveraging **GKE, AKS, OpenShift, and Kubernetes** to deliver highly scalable application platforms.
- Designed and delivered multi-cloud **Platform-as-a-Service (PaaS)** capabilities, enabling engineering teams to consume standardized cloud services through self-service platform patterns.
- Spearheaded the transformation of cloud services into a **self-service PaaS model using the Electrode Framework**, improving developer productivity and accelerating application delivery across global engineering teams.
- Architected **Infrastructure-as-Code and platform automation** using **Terraform, Concord, and enterprise IaC practices**, standardizing provisioning, configuration, and lifecycle management across multi-cloud environments.
- Contributed to the **Terraform AzureRM Provider**, extending provider capabilities for enterprise Azure services and strengthening Azure Infra.
- Integrated **Generative AI and Agentic AI** capabilities into cloud platforms, enabling automation, AI-driven workflows, and scalable AI application patterns.
- Expanded expertise in **GPU-accelerated AI infrastructure and NVIDIA technologies**, including **DGX, CUDA, GPU workloads, and Kubernetes-based AI architectures**, supporting scalable AI/ML workload design.
- Developed cloud-native AI deployment patterns using the **NVIDIA AI software ecosystem**, including **NVIDIA AI Enterprise, NIM, Triton Inference Server, TensorRT/TensorRT-LLM, and GPU Operator**, for scalable model serving and inference architectures.
- Developed and operationalized a multi-tenant GPU infrastructure using Kubernetes and NVIDIA AI Enterprise, supporting scalable AI/ML model training and inference for global engineering teams.
- Enhanced cloud platform security posture by implementing robust DevSecOps practices and integrating automated security scanning into CI/CD pipelines for Walmart's global engineering teams.

Sr Azure Consultant | Tech Mahindra

November 2018 - August 2020, Bengaluru, India

- Led the design and implementation of enterprise **CI/CD and DevOps platforms using Azure DevOps**, supporting automotive engineering workloads and **MLOps lifecycle management**.
- Designed **hybrid and multi-cloud architectures across Azure and Google Cloud**, integrating cloud services with on-premises automotive engineering environments to support connected and autonomous vehicle workloads.
- Developed **GCP-based cloud integration patterns** using **GKE, Cloud Storage, Pub/Sub, BigQuery, and Dataflow**, enabling scalable ingestion, processing, and analytics of automotive and vehicle telemetry data.

- Supported **autonomous and connected-vehicle engineering initiatives for ZF TRW**, integrating cloud-native data and ML workflows for real-time telemetry processing, analytics, and machine-learning-driven insights.
- Integrated **PTC ThingWorx with Azure IoT Hub**, enabling real-time telemetry ingestion from vehicle controllers and supporting predictive analytics across **10,000+ vehicles**.
- Modernized legacy delivery frameworks through **Azure DevOps**, establishing standardized source control, artifact management, and automated build and release pipelines for applications and ML models.
- Engineered scalable MLOps workflows for automotive data, integrating cloud-native ML pipelines with real-time telemetry processing and analytics for ZF TRW.

Sr Development Support Engineer | Xerox Corp Inc

August 2015 - November 2018, Cochin, India

- Supported **Siemens Teamcenter PLM** for Xerox printer product engineering, managing **BOM/product structures, EBOM, engineering data, revisions, and product lifecycle workflows**.
- Developed and customized **Teamcenter solutions using BMIDE and C++**, implementing business rules, data models, extensions, and PLM workflow requirements for printer engineering processes.
- Led enterprise **application migration initiatives to Microsoft Azure**, transitioning on-premises workloads to Azure with a focus on scalability, reliability, and controlled production migration.
- Executed **VMware-to-Azure migrations** and supported **Azure Virtual Machine** workloads as part of enterprise cloud modernization initiatives.
- Modernized legacy applications through **containerization and cloud-native rearchitecture**, transitioning stateful workloads toward stateless services and **Azure App Service Environment (ASE)**.
- Streamlined PLM data management by developing custom Teamcenter reports and dashboards, improving access to engineering data for product lifecycle workflows.

Systems Engineer | Syntel Ltd (Atos Syntel)

June 2012 - August 2015, Pune, India

- Developed and supported enterprise **ASP.NET/.NET applications** across the full SDLC, including design, development, testing, deployment, and production support.
- Built and maintained **IIS-hosted enterprise applications**, focusing on application availability, performance, troubleshooting, and reliable production delivery for insurance workloads.
- Performed **production troubleshooting and root-cause analysis**, delivering application enhancements and managing source control and releases using **TFS and GitHub** within Agile development environments.
- Optimized IIS web server configurations and application pool settings, enhancing response times and stability for critical insurance applications.

EDUCATION

Masters in Information Technology | Sikkim Manipal University

Bachelors of Computer Applications | Calicut University

CERTIFICATIONS

Claude Certified Architect

Azure Solutions Architect Expert

Google Professional Cloud Architect

Azure DevOps Engineer Expert

SKILLS

Cloud & AI Platforms: Microsoft Azure, **Azure Citadel**, Microsoft Foundry, Azure OpenAI, Azure AI Services, Azure Machine Learning, **Core42 Compass**, Google Cloud (GCP), Vertex AI, AWS

AI / Generative AI: Generative AI, Agentic AI, Retrieval-Augmented Generation (RAG), MLOps, AI Platform Engineering, AI Inference, Model Serving

NVIDIA AI: NVIDIA AI Enterprise, NVIDIA NIM, NVIDIA NeMo, Triton Inference Server, TensorRT, TensorRT-LLM, CUDA, GPU Operator, NGC

Cloud & Platform Engineering: Kubernetes, AKS, GKE, Red Hat OpenShift, Docker, Terraform, Infrastructure as Code (IaC), Platform Engineering, Azure App Service Environment

DevOps & Automation: GitHub Actions, Azure DevOps, CI/CD, DevSecOps, Release Management, PowerShell, Concord

Data, Integration & Observability: BigQuery, Dataflow, Azure Storage, Azure SQL, Azure Service Bus, Azure Event Hubs, Azure IoT Hub, Azure Monitor, Log Analytics, Prometheus, Grafana

Programming: Python, Go, C++, Groovy, .NET, ASP.NET

PLM & Engineering: Siemens Teamcenter, BMIDE, BOM / EBOM / MBOM, PTC Windchill, PTC ThingWorx, PTC Integrity
